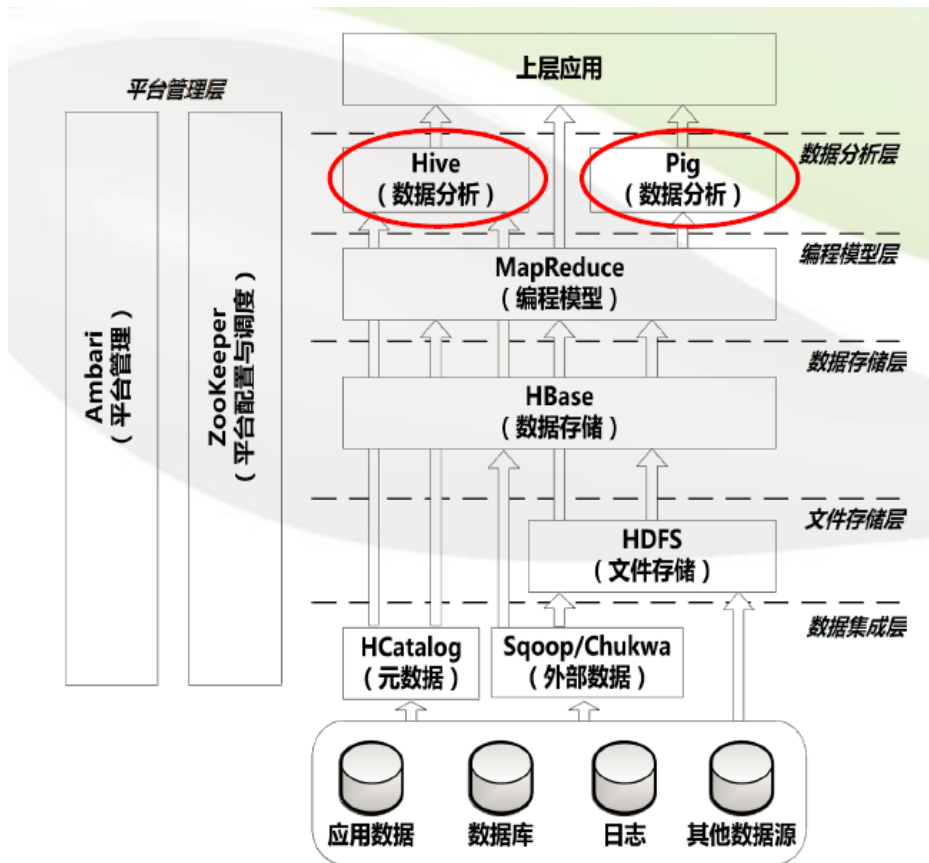


# 议程

- **What is Hive?**
- **Hive Components**
- **What is Hive Data Model?**
- **Underlying Hive Architecture**
- **Using Hive in Practice**

# What is Hive



**Data warehouse infrastructure build on top of Hadoop for querying and managing large data sets**

# Hive存在的必要性

代码

```
#include "mapreduce/mapreduce.h"
```

```
class WordCounter : public Mapper {
```

```
public:
```

```
virtual void Map(const MapInput& input) {
```

```
    const string& text = input.value();
```

```
    const int n = text.size();
```

```
    for (int i = 0; i < n; i++) {
```

```
        while ((i < n) && isspace(text[i])) i++;
```

```
        int start = i;
```

```
        while ((i < n) && !isspace(text[i])) i++;
```

```
        if (start < i)
```

```
            Emit(text.substr(start, i - start, "1");
```

```
    REGISTER_MAPPER(WordCounter);
```

```
class Adder : public Reducer {
```

```
virtual void Reduce(ReduceInput* input) {
```

```
    int64 value = 0;
```

```
    while (!input->done()) {
```

```
        value += StringToInt(input->value());
```

```
        input->NextValue();
```

```
    }
```

```
    Emit(IntToString(value));
```

```
    REGISTER_REDUCER(Adder);
```

```
int main(int argc, char** argv) {
```

```
    ParseCommandLineFlags(argc, argv);
```

```
    MapReduceSpecification spec;
```

```
    for (int i = 1; i < argc; i++) {
```

```
        MapReduceInput* input = spec.add_input();
```

```
        input->set_format("text");
```

```
        input->set_filepattern(argv[i]);
```

```
        input->set_mapper_class("WordCounter");
```

```
    }
```

```
    MapReduceOutput* out = spec.output();
```

```
    out->set_filebase("/gfs/test/freq");
```

```
    out->set_num_tasks(100);
```

```
    out->set_format("text");
```

```
    out->set_reducer_class("Adder");
```

```
    out->set_combiner_class("Adder");
```

```
    spec.set_machines(2000);
```

```
    spec.set_map_megabytes(100);
```

```
    spec.set_reduce_megabytes(100);
```

```
    MapReduceResult result;
```

```
    if (!MapReduce(spec, &result)) abort();
```

```
    return 0;
```

```
}
```

Hive

```
SELECT * FROM log WHERE date > '2012-12-01' ;
```



# Hive的核心设计理念

**A system for managing and querying unstructured data as if it were structured**

- Stores schema in Database
- Uses Map-Reduce for execution
- HDFS for Storage

# Hive能做什么

- **Designed for OLAP**
- **SQL type language for querying**
- **It is familiar, fast, scalable, and extensible**

# Hive不适合做什么


- **Relational database**
- **Designed for Online Transaction Processing (OLTP)**
- **Language for real-time queries and row-level updates**

# Hive的历史

- **Early Hive development work started at Facebook in 2007**
- **Hive is an Apache project under Hadoop**

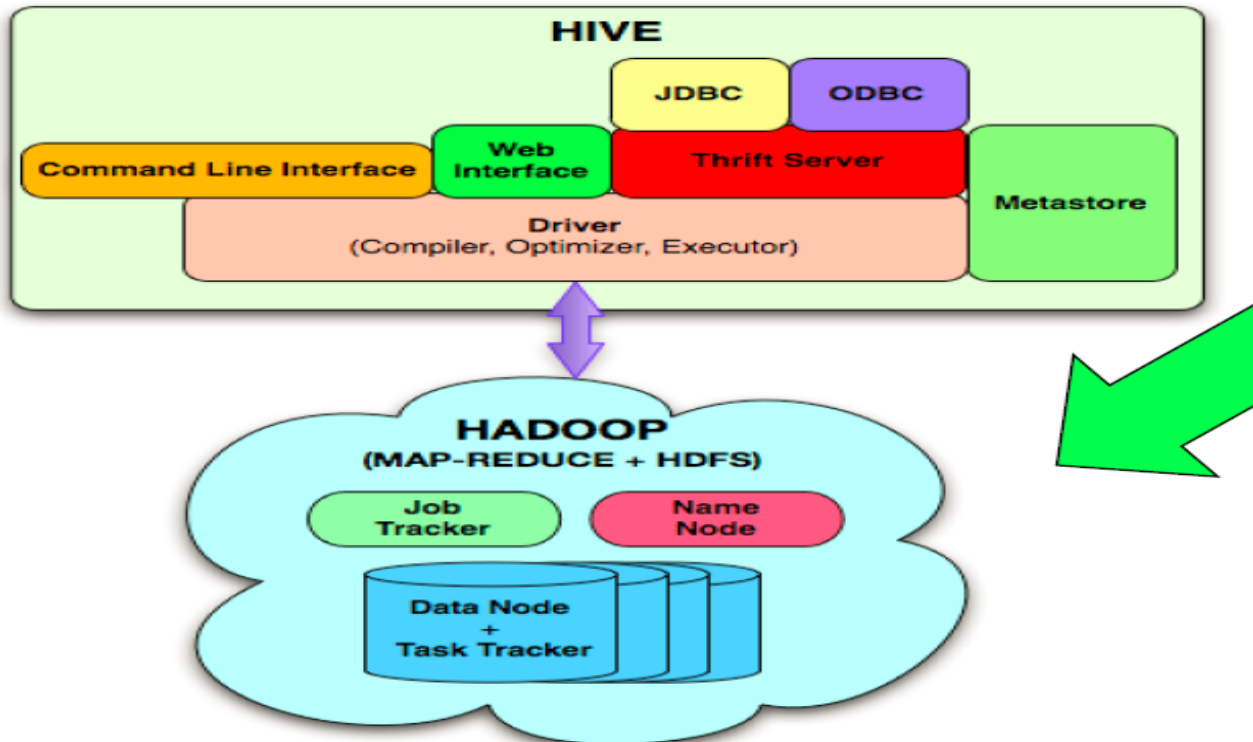
<http://hive.apache.org>

# Hive特性

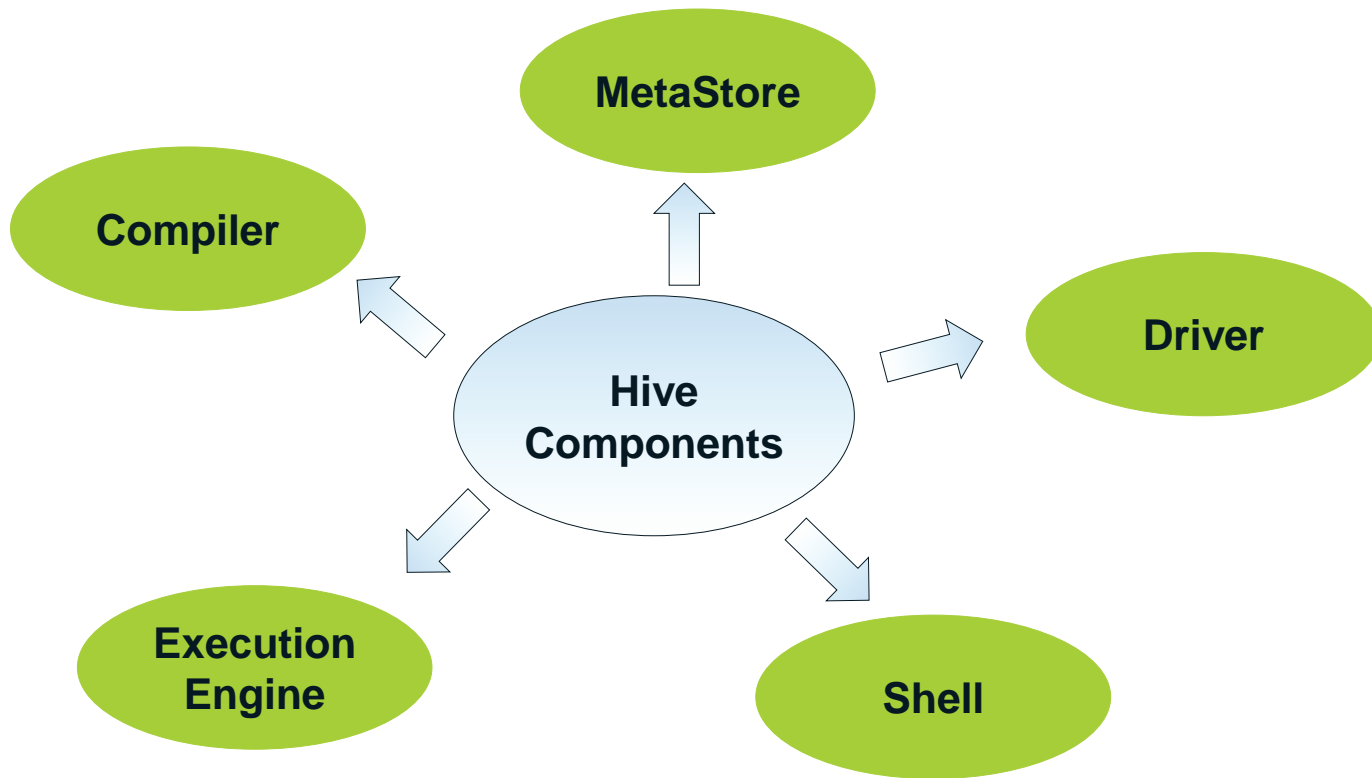




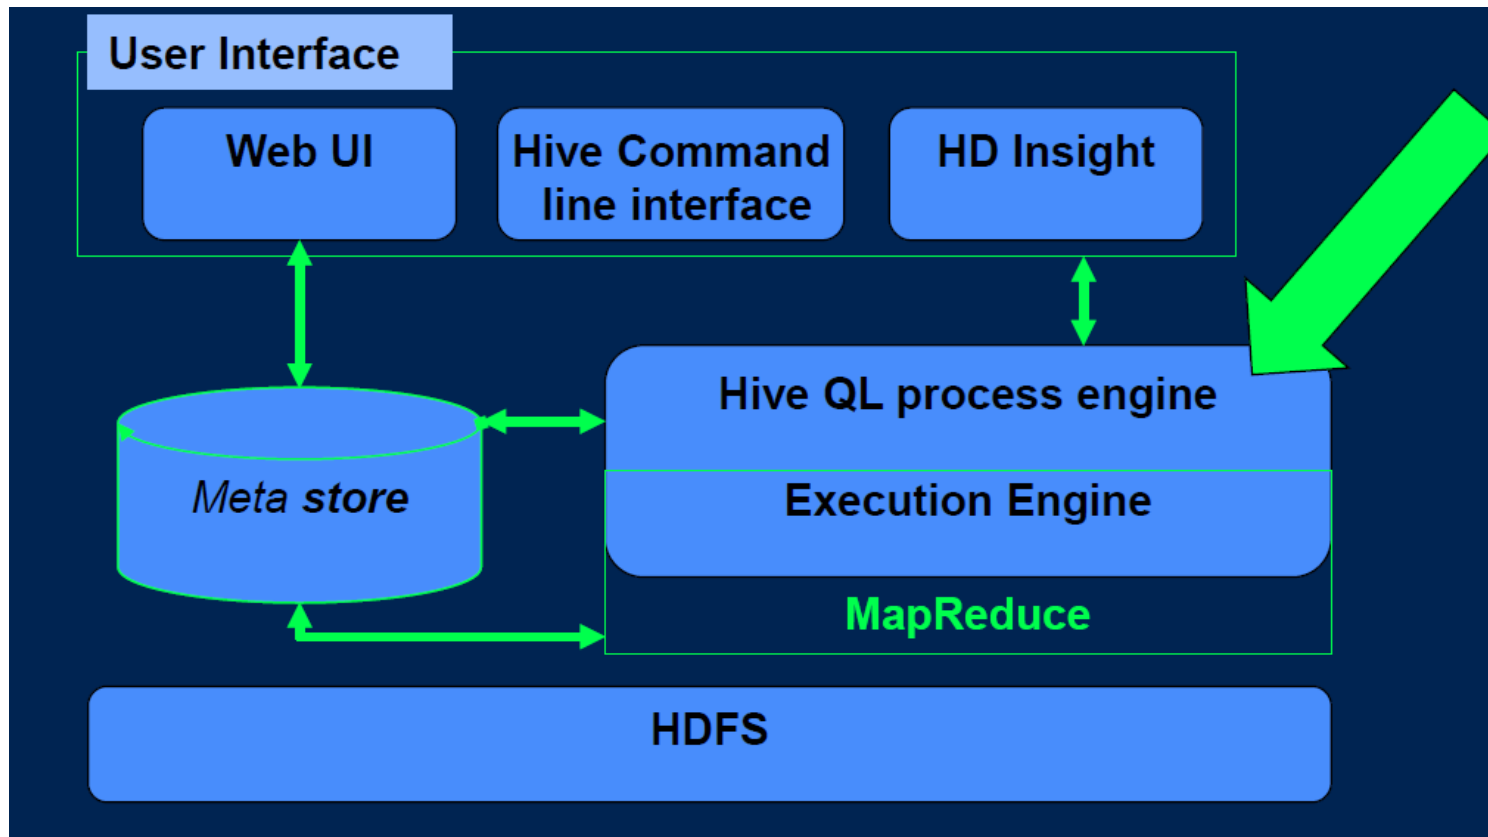
# Hive架构和组件



# Hive组件



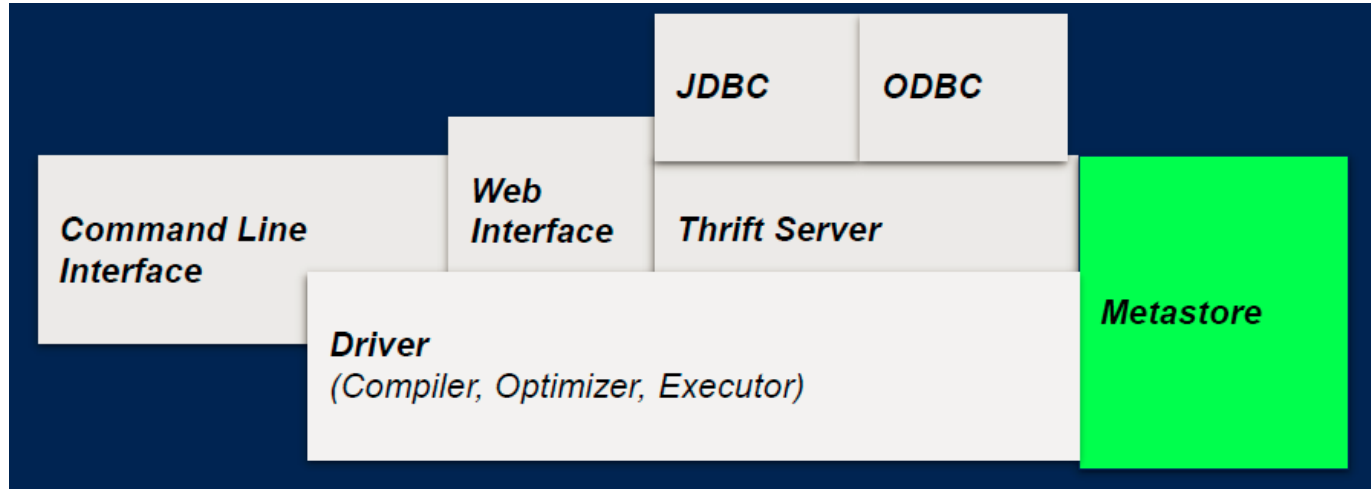
# Hive架构



# Metastore

**Stores the system catalog and meta data about tables, columns, partitions etc.**

**Stored on a traditional RDBMS**



# Driver

**Manages the lifecycle of a HiveQL statement**

**Maintains a session handle and any session statistics**

# Query Compiler

**The component that compiles HiveQL into a directed acyclic graph of map/reduce tasks**

# Optimizer

**Consists of a chain of transformations**

**Performs Column Pruning , Partition Pruning,  
Repartitioning of Data**

# Execution Engine

**Executes the tasks produced by the compiler in proper dependency order**

**Interacts with the underlying Hadoop instance**



# ThriftServer

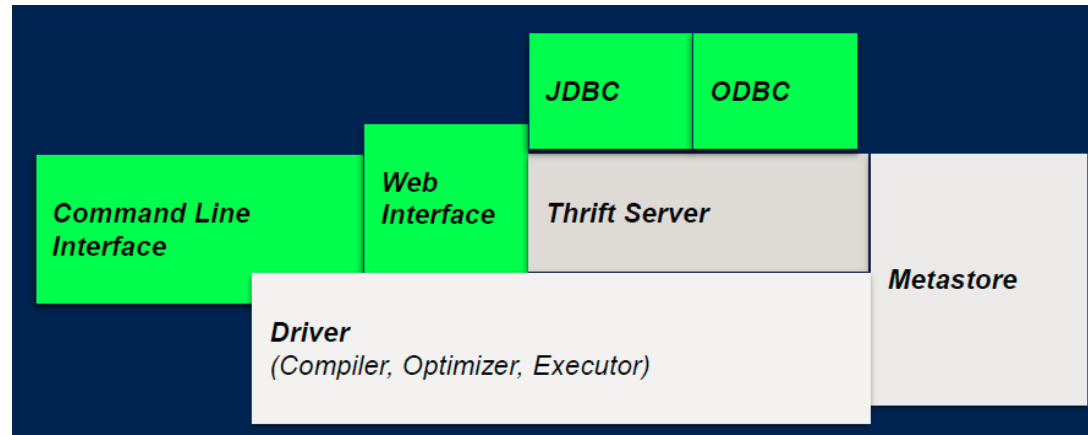
**Provides a Thrift interface and a JDBC/ODBC server Enables Hive integration with other applications**

# Client Components

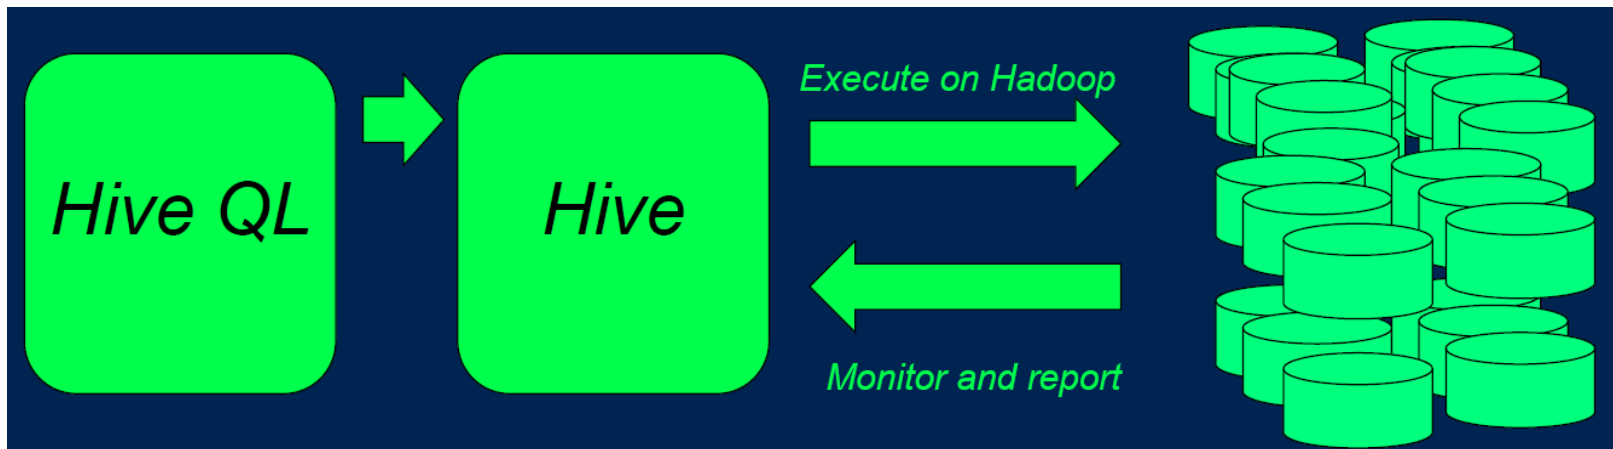
## Command Line Interface(CLI)

## Web UI

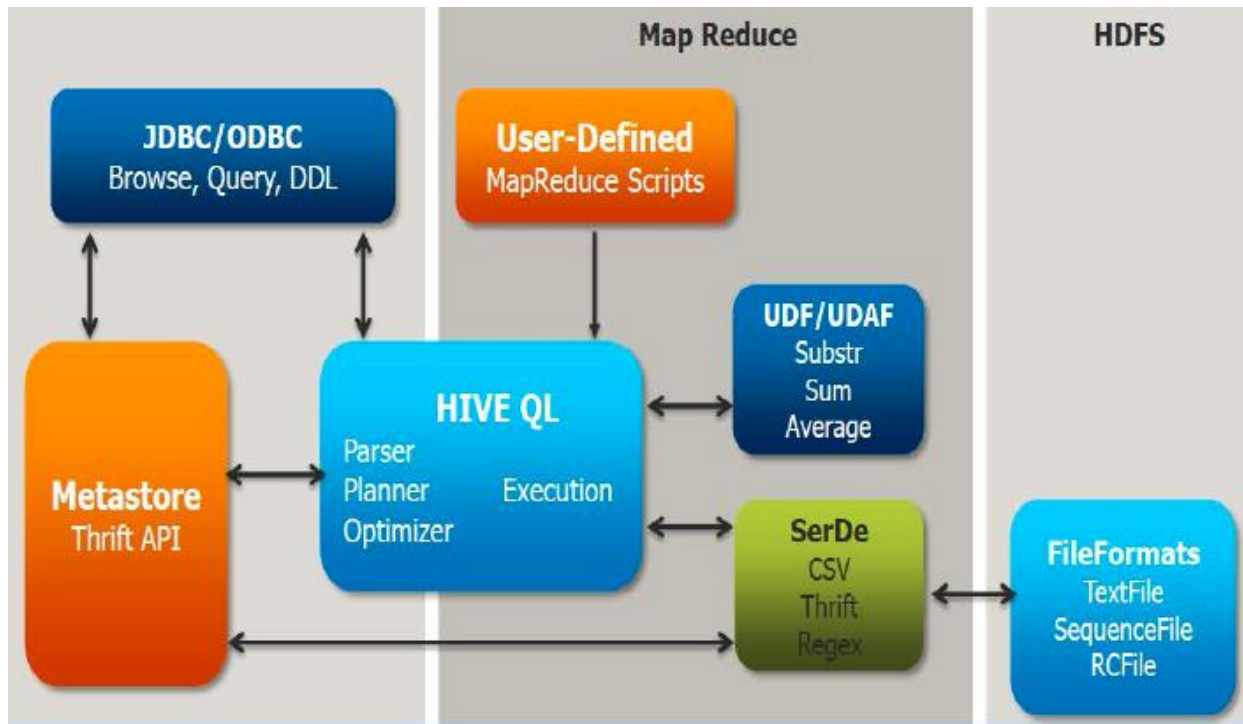
## JDBC/ODBC driver



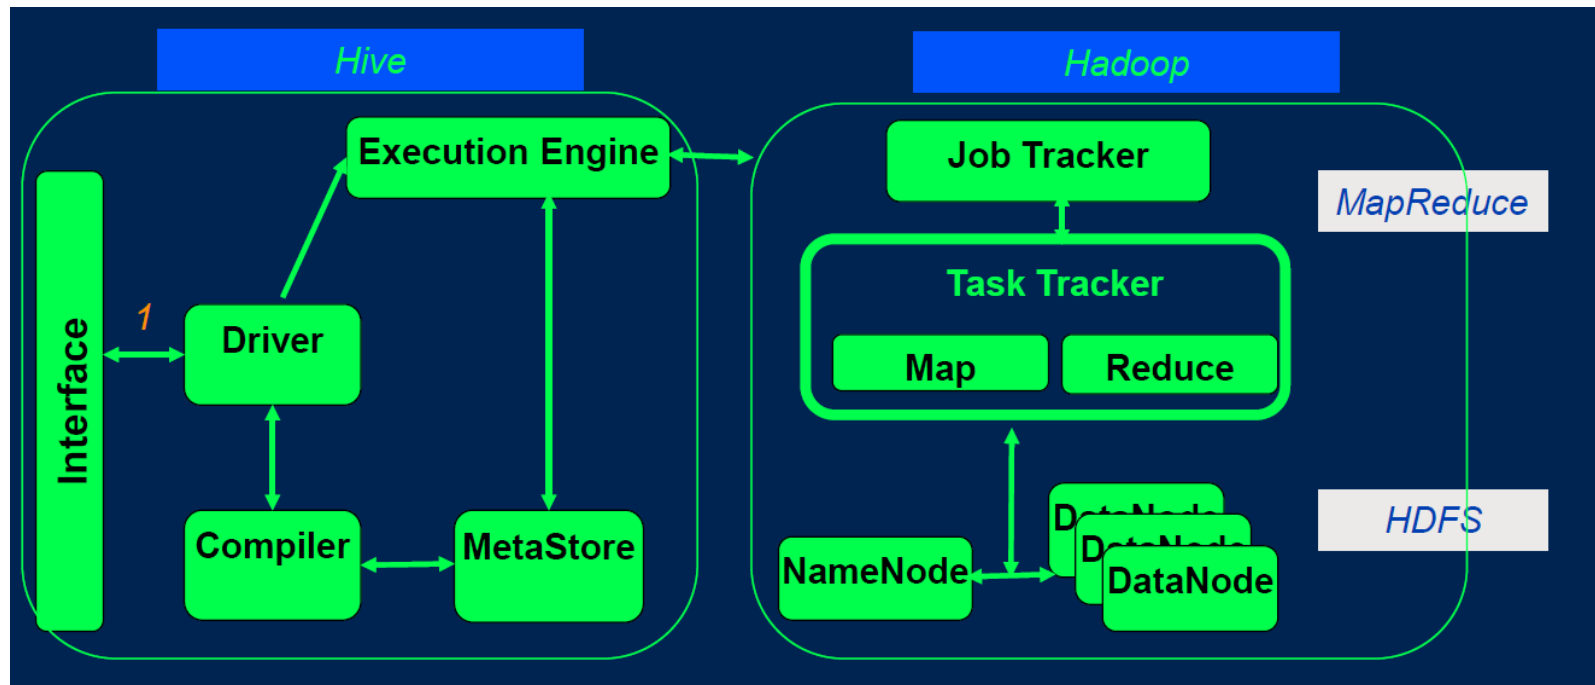
# Hive执行过程



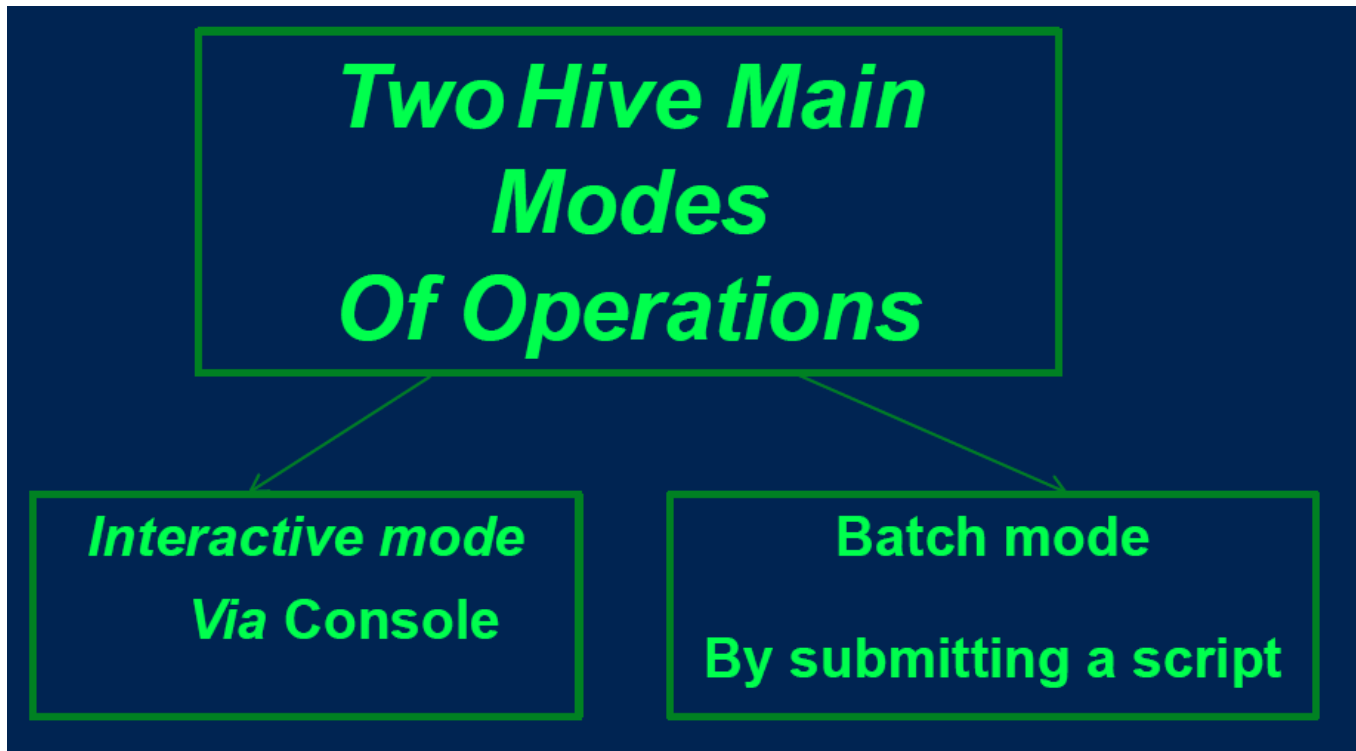
# Hive工作流程



# Hive工作流程



# Hive工作模式



# Hive's Data Units

- Databases
- Tables
- Partitions
- Buckets (or clusters)

*3-Levels: Tables → Partitions → Buckets*

# Data Model

- **Table maps to a HDFS directory**
- **Partition maps to sub-directories under the table**
- **Bucket maps to files under each partition**



# Tables

Similar to tables in relational DBs

Each table has corresponding directory in HDFS

# Partitions

- Analogous to dense indexes on partition columns
- Nested sub-directories in HDFS for each combination of partition column values
- Allows users to efficiently retrieve rows

# Hive Data Structures

- Tables
- Rows
- Columns
- Partitions

# Hive基础数据类型

- Integers
- Floats
- Doubles
- Strings

# Hive扩展数据类型

- Associative arrays : `map<key-type, value-type>`
- Lists : `list<element type>`
- Structs : `struct<file name: file type...>`

# Hive支持的文件类型

**Hive enables users store different file formats**

**Performance improvements**

- TEXTFILE
- SEQUENCEFILE
- ORC
- RCFILE

# Hive Interface

- **Command Line interface**
- **Web interface or Hue**
- **Java Database connectivity**

# Hive Commands

## **Database**

Set of Tables - name conflicts resolution

## **Table**

Set of Rows - have the same columns

## **Row**

A single record - a set of columns

## **Column**

Value and type for a single value



# Tables

- **SHOW TABLES**
- **CREATE TABLE**
- **ALTER TABLE**
- **DROP TABLE**

# Table Commands

**CREATE TABLE** mytable (myint INT, bar STRING)  
**PARTITIONED BY** (ds STRING);

**SHOW TABLES** '.\*my';

**ALTER TABLE** mytable **ADD COLUMNS** (new\_col  
INT);

**DROP TABLE** mytable;

# Hive Query Language

## JOIN

- **SELECT t1.a1 as c1, t2.b1 as c2 FROM t1 JOIN t2 ON (t1.a2 = t2.b2);**

## INSERTION

- **INSERT OVERWRITE TABLE t1 SELECT \* FROM t2;**

# Format rows

**CREATE TABLE mypeople (id INT, name STRING)**

**ROW FORMAT**

**DELIMITED FIELDS TERMINATED BY <output format>**

**LINES TERMINATED BY '\n';**

# Loading data into Hive

## HDFS

- **LOAD DATA INPATH 'mybigdata' [OVERWRITE] INTO TABLE mypeople;**

## Local file system

- **LOAD DATA LOCAL INPATH 'mybigdata' INTO TABLE mypeople;**

## Partitions

- **LOAD DATA INPATH 'myweblogs' INTO TABLE mypeople  
PARTITION (dt=12-12-2020);**

# BUCKETS

**Set hive.enforce.bucketing property to true**

**CREATE TABLE mycustomers(id INT, purchases DOUBLE,  
name STRING)**

**CLUSTERED BY id into 32 BUCKETS;**

**SELECT min(cost) FROM mysales TABLESAMPLE (BUCKET  
10 OUT OF 32 ON rand());**

# VIEWS

## Similar to SQL Views

Virtual table in Metastore

SHOW TABLES

# JOINS

**LEFT OUTER JOIN**

**RIGHT OUTER JOIN**

**FULL OUTER JOIN**

```
hive> SELECT c.ID, c.NAME, c.AGE, o.AMOUNT  
FROM CUSTOMERS c JOIN ORDERS o ON  
(c.ID = o.CUSTOMER_ID);
```



# 外部表

**CREATE EXTERNAL TABLE customers STORED  
AS AVRO**

**LOCATION**

**'hdfs:///user/hive/warehouse/customers'**