

HBase介绍

英特尔软件部
邓 刚

议程

- What is HBase
- HBase Architecture
- HBase Basics

为什么需要NoSQL

传统关系型数据库的优点：

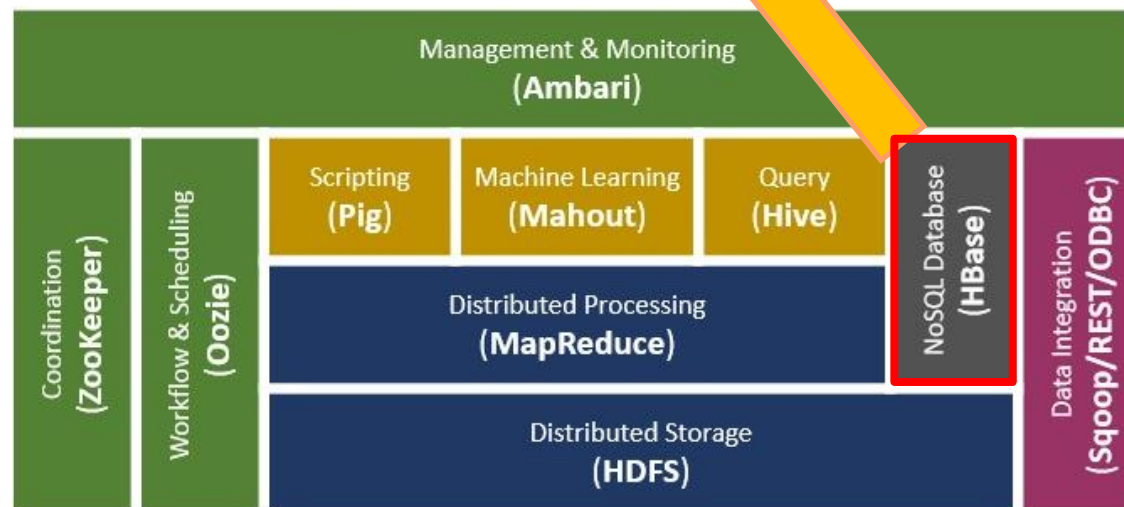
- 严格的一致性模型
- 易于管理和维护
- ◦ ◦

传统关系型数据库的挑战：

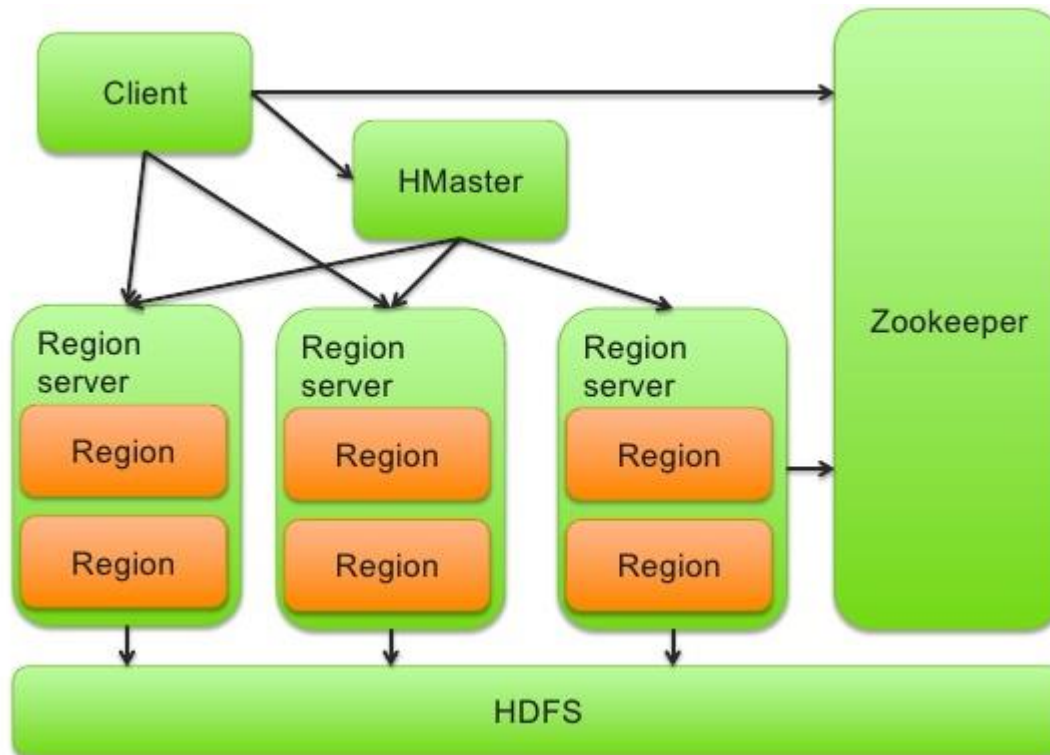
- 扩展瓶颈 - scale up vs. scale out
- 严苛的Schema

What is HBase

- KV 模型
- 分布式 - 可水平扩展到成百上千个节点
- 列式存储 - 灵活的schema, 便于压缩
- 大容量 - PB级数据
- 高性能



HBase Architecture



Hbase Architecture (cont.)

- Table is made up of any number of regions
- Region is specified by its startKey and endKey
 - Empty table: (Table, NULL, NULL)
 - Two-region table: (Table, NULL, "com.cloudera.www") and (Table, "com.cloudera.www", NULL)
- Each region may live on a different node and is made up of several HDFS files and blocks, each of which is replicated by Hadoop

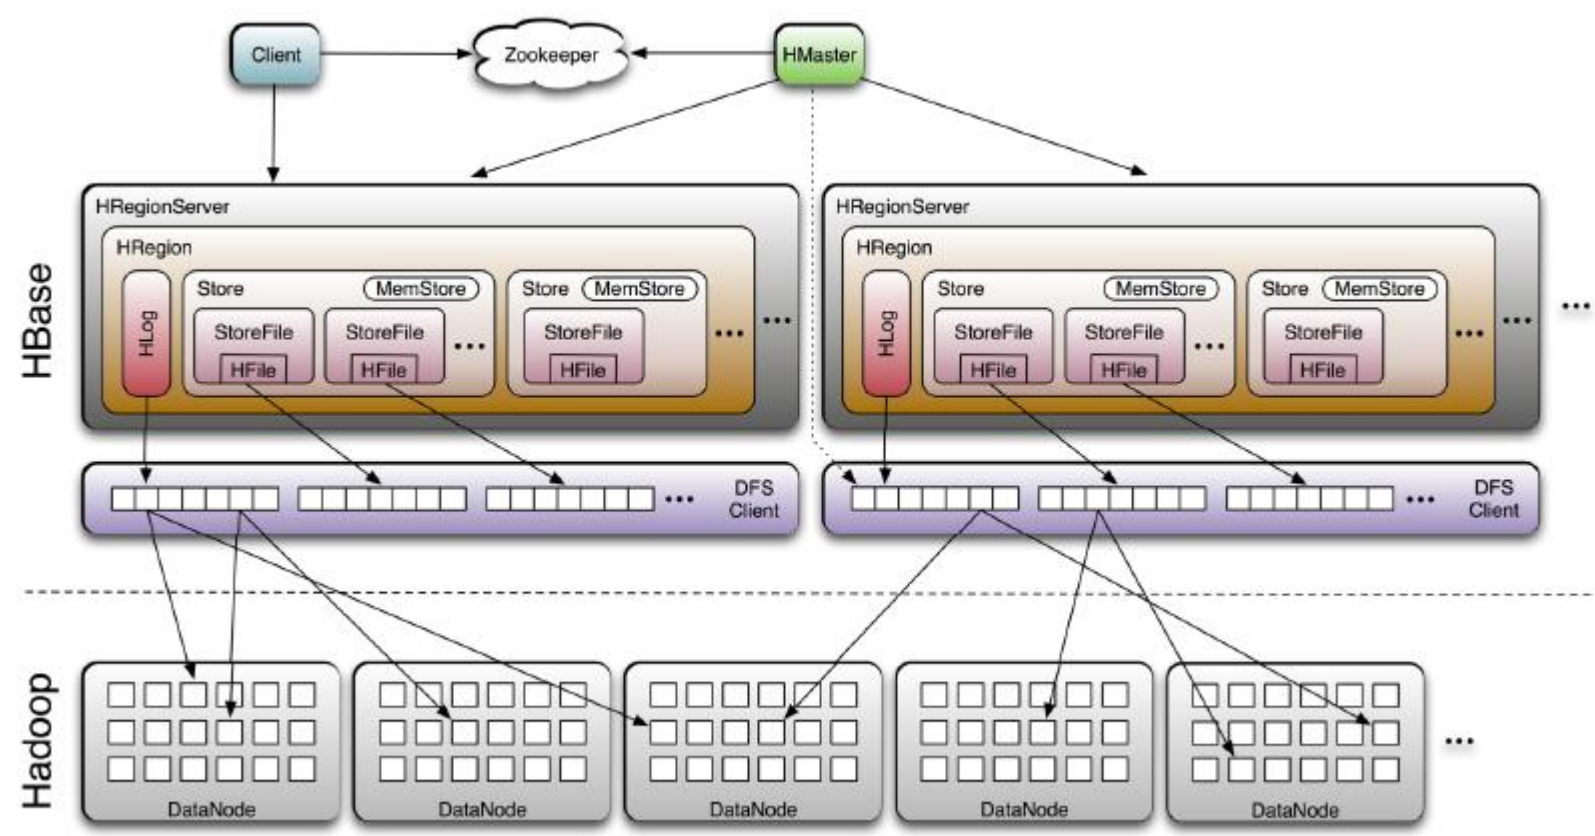
Hbase Architecture (cont.)

- Two types of HBase nodes:

Master and RegionServer

- Special tables -ROOT- and .META. store schema information and region locations
- Master server responsible for RegionServer monitoring as well as assignment and load balancing of regions
- Uses ZooKeeper as its distributed coordination service
 - Manages Master election and server availability

Hbase Architecture (cont.)



Hbase Architecture (cont.)

- Based on Log-Structured Merge-Trees (LSM-Trees)
- Inserts are done in write-ahead log first
- Data is stored in memory and flushed to disk on regular intervals or based on size
- Small flushes are merged in the background to keep number of files small
- Reads read memory stores first and then disk based files second
- Deletes are handled with “tombstone” markers
- Atomicity on row level no matter how many columns
 - keeps locking model easy

HBase Table

	A	B	C	D	E
1					
2					
3					
4					
5					
6					
7					



Row Keys

Column Names, aka Column Qualifiers, aka Column Keys

	A	B	C	D	E
1					
2					
3					
4					
5					
6					
7					



	col-A	col-B	col-Foo	col-XYZ	foobar
row-1					
row-10					
row-18	A18	B18	Foo18	XYZ18	foobar18
row-2					
row-5					
row-6					
row-7					



Row Keys

Column Names, aka Column Qualifiers, aka Column Keys

	col-A	col-B	col-Foo	col-XYZ	foobar
row-1					
row-10					
row-18					
row-2					
row-5					
row-6					
row-7					

HBase Table (cont.)

	col-A	col-B	col-Foo	col-XYZ	foobar
row-1					
row-10					
row-18	A18 - v1	B18 - v3	Foo18 - v1	XYZ18 - v2	foobar18 - v1
row-2					
row-5					
row-6					
row-7					



	Column Family 1		Column Family 2		
	cf1:col-A	cf1:col-B	cf2:col-Foo	cf2:col-XYZ	cf2:foobar
row-1					
row-10					
row-18	A18 - v1	B18 - v3	Foo18 - v1	XYZ18 - v2	foobar18 - v1
row-2		Peter - v2 Bob - v1		Mary - v1	
row-5					
row-6					
row-7					

Coordinates for a Cell: Row Key → Column Family Name → Column Qualifier → Version



	Column Family 1		Column Family 2		
	cf1:col-A	cf1:col-B	cf2:col-Foo	cf2:col-XYZ	cf2:foobar
row-1					
row-10					
row-18	A18 - v1	B18 - v3	Foo18 - v1	XYZ18 - v2	foobar18 - v1
row-2					
row-5					
row-6					
row-7					

Physical Coordinates for a Cell: Region Directory → Column Family Directory
→ Row Key → Column Family Name → Column Qualifier → Version

HBase Table (cont.)

- Tables are sorted by the *Row Key* in lexicographical order
- Table schema only defines its *Column Families*
 - Each family consists of any number of *Columns*
 - Each column consists of any number of *Versions*
- Columns only exist when inserted, NULLs are free
- Columns within a family are sorted and stored together
- Everything except table names are byte[]

(Table, Row, Family:Column, Timestamp) → Value

MemStores

- After data is written to the WAL the RegionServer saves KeyValues in **memory store**
- Flush to disk based on size, see *hbase.hregion.memstore.flush.size*
- Default size is **64MB**
- Uses **snapshot** mechanism to write flush to disk while still serving from it and accepting new data at the same time
- Snapshots are released when flush has succeeded

Block Cache

- Acts as very large, in-memory **distributed** cache
- Assigned a large part of the JVM **heap** in the RegionServer process, see *hfile.block.cache.size*
- Optimizes **reads** on subsequent columns and rows
- Has **priority** to keep “in-memory” column families in cache

```
if(inMemory) {  
    this.priority = BlockPriority.MEMORY;  
} else {  
    this.priority = BlockPriority.SINGLE;  
}
```

- Cache needs to be used properly to get best read performance
 - Turn off block cache on operations that cause large churn
 - Store related data “close” to each other
- Uses **LRU** cache with threaded (asynchronous) evictions based on priorities

Block Cache

- **General Concepts**
 - Two types: **Minor** and **Major** Compactions
 - Asynchronous and transparent to client
 - Manage file bloat from MemStore flushes
- **Minor Compactions**
 - Combine last “few” flushes
 - Triggered by number of storage files
- **Major Compactions**
 - Rewrite **all** storage files
 - Drop deleted data and those values exceeding TTL and/or number of versions
 - Triggered by time threshold
 - Cannot be scheduled automatically starting at a specific time (bummer!)
 - May (most definitely) tax overall HDFS IO performance


Tip: Disable major compactions and schedule to run manually (e.g. cron) at off-peak times

Bloom Filters

- Defines a filter that allows to determine if a store file does **not** contain a row or column
- Error rate can control overhead but is usually very low, 1% or less
- Stored with each storage file on flush and compactions
- Good for large regions with many distinct row keys and many expected misses

Key Design

- Based on access pattern, either use sequential or random keys
- Often a combination of both is needed
 - Overcome architectural limitations
- Neither is necessarily bad
 - Use bulk import for sequential keys and reads
 - Random keys are good for random access patterns



ColumnFamily vs. Column

- Use only a few column families
 - Causes many files that need to stay open per region plus class overhead per family
- Best used when logical separation between data and meta columns
- Sorting per family can be used to convey application logic or access pattern
- Define compression or in-memory attributes to optimize access and performance

Web Crawl Example

- Canonical use-case for BigTable
- Store web crawl data
 - Table **webtable** with family **content** and **meta**
 - Row is reversed URL with Columns
 - *content:data* stores the raw crawled data
 - *meta:language* stores http language header
 - *meta:type* stores http content-type header
 - While processing raw data for hyperlinks and images, add families **links** and **images**
 - *links:<rurl>* column for each hyperlink
 - *images:<rurl>* column for each image