


Spark介绍

为什么需要Spark

- Hadoop write intermediate data into HDFS



- Map/Reduce task launch overhead, so do Hive
 - Task launch delay 5 ~ 10 seconds
- No analysis function in Hbase, without Hive
- Low memory and CPU utilization, I/O is the bottle neck
- Can not support the ad-hoc queries

Spark与Hadoop MapReduce的对比

- Spark takes the concepts of MapReduce to the next level
 - Higher level API = faster, easier development
 - Low latency = near real-time processing
 - In-memory data storage = up to 100x performance improvement



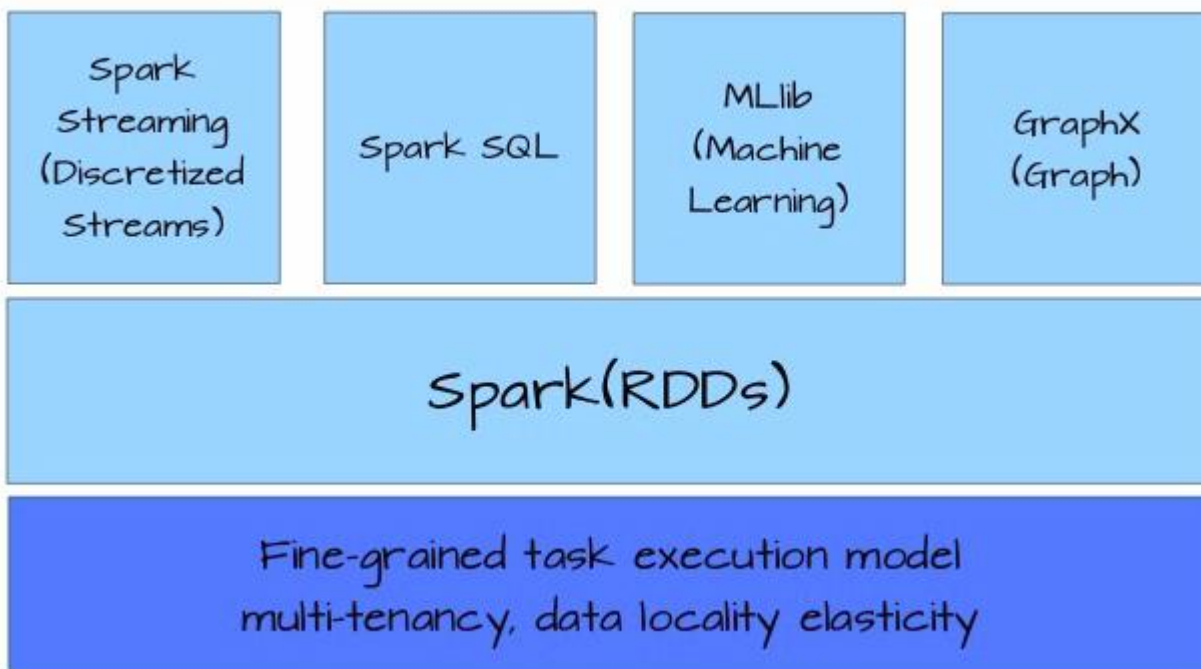
```
sc.textFile(file) \  
  .flatMap(lambda s: s.split()) \  
  .map(lambda w: (w,1)) \  
  .reduceByKey(lambda v1,v2: v1+v2) \  
  .saveAsTextFile(output)
```



```
public class WordCount {  
    public static void main(String[] args) throws  
    {  
        Job job = new Job();  
        job.setJarByClass(WordCount.class);  
        job.setJobName("Word Count");  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        boolean success = job.waitForCompletion(true);  
        System.exit(success ? 0 : 1);  
    }  
}  
  
public class WordMapper extends Mapper<LongWritable, Text, Text,  
    IntWritable> {  
    public void map(LongWritable key, Text value,  
        Context context) throws IOException, InterruptedException {  
        String line = value.toString();  
        for (String word : line.split("\\W+")) {  
            if (word.length() > 0) {  
                context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
}  
  
public class SumReducer extends Reducer<Text, IntWritable, Text,  
    IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable>  
        values, Context context) throws IOException, InterruptedException {  
        int wordCount = 0;  
        for (IntWritable value : values) {  
            wordCount += value.get();  
        }  
        context.write(key, new IntWritable(wordCount));  
    }  
}
```




Spark的软件栈



支持的语言: Java, Python, Scala

内存对Spark的影响




Performance without Memory

	Hadoop Record	Spark 100TB	Spark 1PB
Data Size	102.5TB	100TB	1000TB
Time	72 min	23 min	234 min
# Cores	50400	6592	6080
Rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Environment	Dedicate	Cloud (EC2)	Cloud (EC2)

RDD (Resilient Distributed Datasets)

特点

- 容错性
- 并行
- 迭代操作

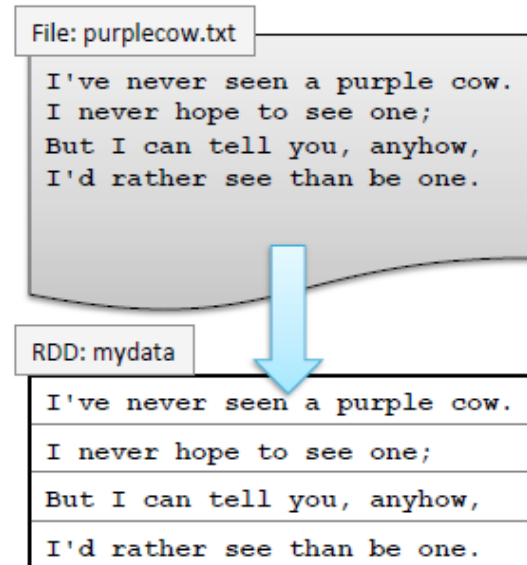


生成RDD

- 从文件生成
- 从来自内存的数据生成
- 从其它RDD

File-Based RDDs


- For file-based RDDs, use `SparkContext.textFile`
 - Accepts a single file, a wildcard list of files, or a comma-separated list of files
 - Examples
 - `sc.textFile("myfile.txt")`
 - `sc.textFile("mydata/*.log")`
 - `sc.textFile("myfile1.txt,myfile2.txt")`
 - Each line in the file(s) is a separate record in the RDD




RDD Operations

- Two types of RDD operations

- Actions – return values




- Transformations – define a new RDD based on the current one(s)



RDD Operations: Actions

- Some common actions

- `count()` – return the number of elements
- `take(n)` – return an array of the first n elements
- `collect()` – return an array of all elements
- `saveAsTextFile(file)` – save to text file(s)



```
> mydata =  
  sc.textFile("purplecow.txt")  
  
> mydata.count()  
4  
  
> for line in mydata.take(2):  
  print line  
I've never seen a purple cow.  
I never hope to see one;
```

```
> val mydata =  
  sc.textFile("purplecow.txt")  
  
> mydata.count()  
4  
  
> for (line <- mydata.take(2))  
  println(line)  
I've never seen a purple cow.  
I never hope to see one;
```

RDD Operations: Operations


- Transformations create a new RDD from an existing one

- RDDs are immutable

- Data in an RDD is never changed
- Transform in sequence to modify the data as needed

- Some common transformations

- `map(function)` – creates a new RDD by performing a function on each record in the base RDD
- `filter(function)` – creates a new RDD by including or excluding each record in the base RDD according to a boolean function



Spark SQL: Relational Data Processing in Spark

- **Integrated** – Seamlessly mix SQL queries with Spark programs.
- **Unified Data Access** – Load and query data from a variety of sources.
- **Hive Compatibility** – Run unmodified Hive queries on existing warehouses.
- **Standard Connectivity** – Connect through JDBC or ODBC.

Language API



Data Sources

Query on Object

```
val sqlContext = new org.apache.spark.sql.SQLContext(sc)
import sqlContext._

// Define the schema using a case class.
case class Person(name: String, age: Int)

// Create an RDD of Person objects and register it as a table.
val people = sc.textFile("examples/src/main/resources/
people.txt").map(_._split(",")).map(p => Person(p(0), p(1).trim.toInt))



people.registerAsTable("people")

// SQL statements can be run by using the sql methods provided by sqlContext.
val teenagers = sql("SELECT name FROM people WHERE age >= 13 AND age <= 19")
```



Spark SQL – Query in HiveQL

```
//val sc: SparkContext // An existing SparkContext.  
//NB: example on laptop lacks a Hive MetaStore  
val hiveContext = new org.apache.spark.sql.hive.HiveContext(sc)  
  
// Importing the SQL context gives access to all the  
// public SQL functions and implicit conversions.  
import hiveContext._  
  
hql("CREATE TABLE IF NOT EXISTS src (key INT, value STRING)")  
hql("LOAD DATA LOCAL INPATH 'examples/src/main/resources/kv1.txt' INTO TABLE src")  
  
// Queries are expressed in HiveQL  
hql("FROM src SELECT key, value").collect().foreach(println)
```

Spark SQL Performance




110GB of data
after columnar
compression
with Parquet



Spark Streaming

Run a streaming computation as a **series of very small, deterministic batch jobs**

- Batch sizes as low as ½ second, latency of about 1 second
- Potential for combining batch processing and streaming processing in the same system



Spark Streaming Example

```
import org.apache.spark.streaming._
import org.apache.spark.streaming.StreamingContext._

// Create a StreamingContext with a SparkConf configuration
val ssc = new StreamingContext(sparkConf, Seconds(10))

// Create a DStream that will connect to serverIP:serverPort
val lines = ssc.socketTextStream(serverIP, serverPort)

// Split each line into words
val words = lines.flatMap(_.split(" "))

// Count each word in each batch
val pairs = words.map(word => (word, 1))
val wordCounts = pairs.reduceByKey(_ + _)
```


Comparison with Storm and S4

Higher throughput than Storm

- Spark Streaming: **670k** records/second/node
- Storm: **115k** records/second/node
- Apache S4: 7.5k records/second/node

